

Big Data Programming Assignment 5

Susanth Dasari

Big Data Programming class - Assignment 5

Susanth Dasari

```
library(caret)
library(gbm)
library(RANN)
library(ggplot2)
```

Using R 3.5.1 (or later) and caret's Animal Scat Data dataset (Hint: data(scats))

```
data(scats)
```

```
str(scats)
```

```
> data(scats)
> str(scats)
'data.frame': 110 obs. of 19 variables:
 $ Species : Factor w/ 3 levels "bobcat","coyote",...: 2 2 1 2 2 2 1 1 1 1 ...
 $ Month   : Factor w/ 9 levels "April","August",...: 4 4 4 4 4 4 4 4 4 4 ...
 $ Year    : int 2012 2012 2012 2012 2012 2012 2012 2012 2012 2012 ...
 $ Site    : Factor w/ 2 levels "ANNU","YOLA": 2 2 2 2 2 2 1 1 1 1 ...
 $ Location: Factor w/ 3 levels "edge","middle",...: 1 1 2 2 1 1 3 3 3 2 ...
 $ Age     : int 5 3 3 5 5 5 1 3 5 5 ...
 $ Number  : int 2 2 2 2 4 3 5 7 2 1 ...
 $ Length  : num 9.5 14 9 8.5 8 9 6 5.5 11 20.5 ...
 $ Diameter: num 25.7 25.4 18.8 18.1 20.7 21.2 15.7 21.9 17.5 18 ...
 $ Taper   : num 41.9 37.1 16.5 24.7 20.1 28.5 8.2 19.3 29.1 21.4 ...
 $ TI      : num 1.63 1.46 0.88 1.36 0.97 1.34 0.52 0.88 1.66 1.19 ...
 $ Mass    : num 15.9 17.6 8.4 7.4 25.4 ...
 $ d13C    : num -26.9 -29.6 -28.7 -20.1 -23.2 ...
 $ d15N    : num 6.94 9.87 8.52 5.79 7.01 8.28 4.2 3.89 7.34 6.06 ...
 $ CN      : num 8.5 11.3 8.1 11.5 10.6 9 5.4 5.6 5.8 7.7 ...
 $ ropey   : int 0 0 1 1 0 1 1 0 0 1 ...
 $ segmented: int 0 0 1 0 1 0 1 1 1 1 ...
 $ flat    : int 0 0 0 0 0 0 0 0 0 0 ...
 $ scrape  : int 0 0 1 0 0 0 1 0 0 0 ...
```

1

Set the Species column as the target/outcome and convert it to numeric. (5 points)

```
target <- ifelse(scats$Species=="bobcat",0,ifelse(scats$Species=="coyote",1,2))
```

```
str(target)
```

```
> ### 1
> ### Set the Species column as the target/outcome and convert it to numeric. (5 points)
> target <- ifelse(scats$Species=="bobcat",0,ifelse(scats$Species=="coyote",1,2))
> str(target)
 num [1:110] 1 1 0 1 1 1 0 0 0 0 ...
> |
```

2

Remove the Month, Year, Site, Location features. (5 points)

```
scats$Month <- NULL
```

```
scats$Year <- NULL
```

```
scats$Site <- NULL
```

```
scat$Location <- NULL
```

```
> ### 2
> ### Remove the Month, Year, Site, Location features. (5 points)
> scat$Month <- NULL
> scat$Year <- NULL
> scat$Site <- NULL
> scat$Location <- NULL
> str(scat)
'data.frame': 110 obs. of 15 variables:
 $ Species : Factor w/ 3 levels "bobcat","coyote",...: 2 2 1 2 2 2 1 1 1 1 ...
 $ Age     : int  5 3 3 5 5 5 1 3 5 5 ...
 $ Number  : int  2 2 2 2 4 3 5 7 2 1 ...
 $ Length  : num  9.5 14 9 8.5 8 9 6 5.5 11 20.5 ...
 $ Diameter: num  25.7 25.4 18.8 18.1 20.7 21.2 15.7 21.9 17.5 18 ...
 $ Taper   : num  41.9 37.1 16.5 24.7 20.1 28.5 8.2 19.3 29.1 21.4 ...
 $ TI      : num  1.63 1.46 0.88 1.36 0.97 1.34 0.52 0.88 1.66 1.19 ...
 $ Mass    : num  15.9 17.6 8.4 7.4 25.4 ...
 $ d13C    : num  -26.9 -29.6 -28.7 -20.1 -23.2 ...
 $ d15N    : num  6.94 9.87 8.52 5.79 7.01 8.28 4.2 3.89 7.34 6.06 ...
 $ CN      : num  8.5 11.3 8.1 11.5 10.6 9 5.4 5.6 5.8 7.7 ...
 $ ropey   : int  0 0 1 1 0 1 1 0 0 1 ...
 $ segmented: int  0 0 1 0 1 0 1 1 1 1 ...
 $ flat    : int  0 0 0 0 0 0 0 0 0 0 ...
 $ scrape  : int  0 0 1 0 0 0 1 0 0 0 ...
```

```
### 3
```

Check if any values are null. If there are, impute missing values using KNN. (10 points)

```
sum(is.na(scat))
```

```
> ### 3
> ### Check if any values are null. If there are, impute missing values using KNN. (10 points)
> sum(is.na(scat))
[1] 47
```

#Imputing missing values using KNN.Also centering and scaling numerical columns

```
preProcValues <- preProcess(scat, method = "knnImpute")
```

```
scat_processed <- predict(preProcValues, scat)
```

```
sum(is.na(scat_processed))
```

```
> #Imputing missing values using KNN.Also centering and scaling numerical columns
> preProcValues <- preProcess(scat, method = "knnImpute")
> scat_processed <- predict(preProcValues, scat)
> sum(is.na(scat_processed))
[1] 0
```

4

Converting every categorical variable to numerical (if needed). (5 points)

```
str(scat_processed)
```

```
> ### 4
> ### Converting every categorical variable to numerical (if needed). (5 points)
> str(scat_processed)
'data.frame': 110 obs. of 15 variables:
 $ Species : Factor w/ 3 levels "bobcat","coyote",...: 2 2 1 2 2 2 1 1 1 1 ...
 $ Age      : num 1.207 -0.252 -0.252 1.207 1.207 ...
 $ Number   : num -0.433 -0.433 -0.433 -0.433 0.968 ...
 $ Length   : num 0.0587 1.3679 -0.0867 -0.2322 -0.3777 ...
 $ Diameter : num 1.8396 1.7623 0.0622 -0.1181 0.5516 ...
 $ Taper     : num 0.961 0.642 -0.726 -0.182 -0.487 ...
 $ TI        : num 0.0283 -0.1406 -0.7171 -0.24 -0.6277 ...
 $ Mass      : num 0.388 0.583 -0.458 -0.571 1.469 ...
 $ d13C      : num 0.00468 -1.26856 -0.85947 3.12113 1.66403 ...
 $ d15N      : num -0.165 0.807 0.359 -0.546 -0.141 ...
 $ CN        : num 0.0276 0.7922 -0.0816 0.8468 0.6011 ...
 $ ropey     : num -1.131 -1.131 0.876 0.876 -1.131 ...
 $ segmented: num -1.131 -1.131 0.876 -1.131 0.876 ...
 $ flat      : num -0.239 -0.239 -0.239 -0.239 -0.239 ...
 $ scrape    : num -0.217 -0.217 4.562 -0.217 -0.217 ...
```

There are no categorical variables remaining

#Converting the dependent variable back to numeric categorical

```
scat_processed$Species<-as.factor(target)
```

```
str(scat_processed)
```

```
> #Converting the dependent variable back to numeric categorical
> scat_processed$Species<-as.factor(target)
> str(scat_processed)
'data.frame': 110 obs. of 15 variables:
 $ Species : Factor w/ 3 levels "0","1","2": 2 2 1 2 2 2 1 1 1 1 ...
 $ Age      : num 1.207 -0.252 -0.252 1.207 1.207 ...
 $ Number   : num -0.433 -0.433 -0.433 -0.433 0.968 ...
 $ Length   : num 0.0587 1.3679 -0.0867 -0.2322 -0.3777 ...
 $ Diameter : num 1.8396 1.7623 0.0622 -0.1181 0.5516 ...
 $ Taper     : num 0.961 0.642 -0.726 -0.182 -0.487 ...
 $ TI        : num 0.0283 -0.1406 -0.7171 -0.24 -0.6277 ...
 $ Mass      : num 0.388 0.583 -0.458 -0.571 1.469 ...
 $ d13C      : num 0.00468 -1.26856 -0.85947 3.12113 1.66403 ...
 $ d15N      : num -0.165 0.807 0.359 -0.546 -0.141 ...
 $ CN        : num 0.0276 0.7922 -0.0816 0.8468 0.6011 ...
 $ ropey     : num -1.131 -1.131 0.876 0.876 -1.131 ...
 $ segmented: num -1.131 -1.131 0.876 -1.131 0.876 ...
 $ flat      : num -0.239 -0.239 -0.239 -0.239 -0.239 ...
 $ scrape    : num -0.217 -0.217 4.562 -0.217 -0.217 ...
```

5

With a seed of 100, 75% training, 25% testing.

Build the following models: randomforest, neural net, naive bayes and GBM.

For these models display

a) model summarization and

b) plot variable of importance, for the predictions (use the prediction set) display

c) confusion matrix (60 points)

Splitting Data Using Caret

#Splitting training set into two parts based on outcome: 75% and 25%

```

set.seed(100)
index <- createDataPartition(scat_processed$Species, p=0.75, list=FALSE)
trainSet <- scat_processed[ index,]
testSet <- scat_processed[-index,]

outcomeName<-'Species'
predictors<-names(trainSet)[!names(trainSet) %in% outcomeName]

##### Training Models Using Caret #####
model_rf<-train(trainSet[,predictors],trainSet[,outcomeName],method='rf')
model_nnet<-train(trainSet[,predictors],trainSet[,outcomeName],method='nnet')
model_nb<-train(trainSet[,predictors],trainSet[,outcomeName],method='naive_bayes')
model_gbm<-train(trainSet[,predictors],trainSet[,outcomeName],method='gbm')

```

Model Statistics for Random Forrest

```
print(model_rf)
```

```

> print(model_rf)
Random Forest

83 samples
14 predictors
3 classes: '0', '1', '2'

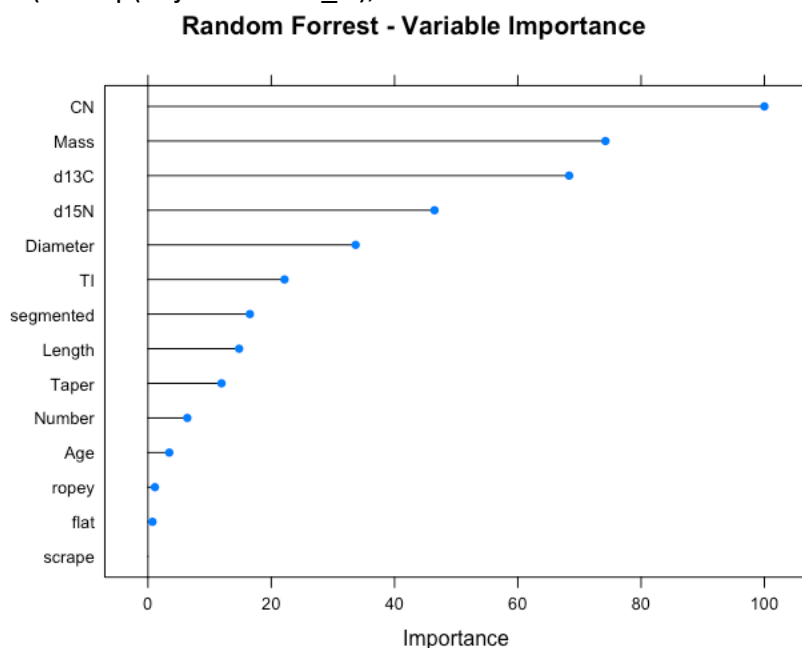
No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
Resampling results across tuning parameters:

mtry  Accuracy  Kappa
  2    0.6291812 0.3672963
  8    0.6493452 0.4182099
 14    0.6507490 0.4248715

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was mtry = 14.
> |

```

```
plot(varImp(object=model_rf),main="Random Forrest - Variable Importance")
```



```
predictions<-predict.train(object=model_rf,testSet[,predictors],type="raw")
confusionMatrix(predictions,testSet[,outcomeName])
```

```
> predictions<-predict.train(object=model_rf,testSet[,predictors],type="raw")
> confusionMatrix(predictions,testSet[,outcomeName])
```

Confusion Matrix and Statistics

	Reference		
Prediction	0	1	2
0	14	2	2
1	0	5	0
2	0	0	4

Overall Statistics

Accuracy : 0.8519
 95% CI : (0.6627, 0.9581)
 No Information Rate : 0.5185
 P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7416

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.6923	1.0000	1.0000
Pos Pred Value	0.7778	1.0000	1.0000
Neg Pred Value	1.0000	0.9091	0.9130
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.6667	0.1852	0.1481
Balanced Accuracy	0.8462	0.8571	0.8333

> |

Model Statistics for Neural Networks

```
print(model_nnet)
```

```
> print(model_nnet)
```

Neural Network

83 samples

14 predictors

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

size	decay	Accuracy	Kappa
1	0e+00	0.5843441	0.3169681
1	1e-04	0.5408497	0.2553937
1	1e-01	0.6296692	0.3696320
3	0e+00	0.6424186	0.4189644
3	1e-04	0.6615715	0.4482782
3	1e-01	0.6908921	0.4788325
5	0e+00	0.6371559	0.4079960
5	1e-04	0.6407128	0.4052624
5	1e-01	0.6906994	0.4805664

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were size = 3 and decay = 0.1.

```
> |
```

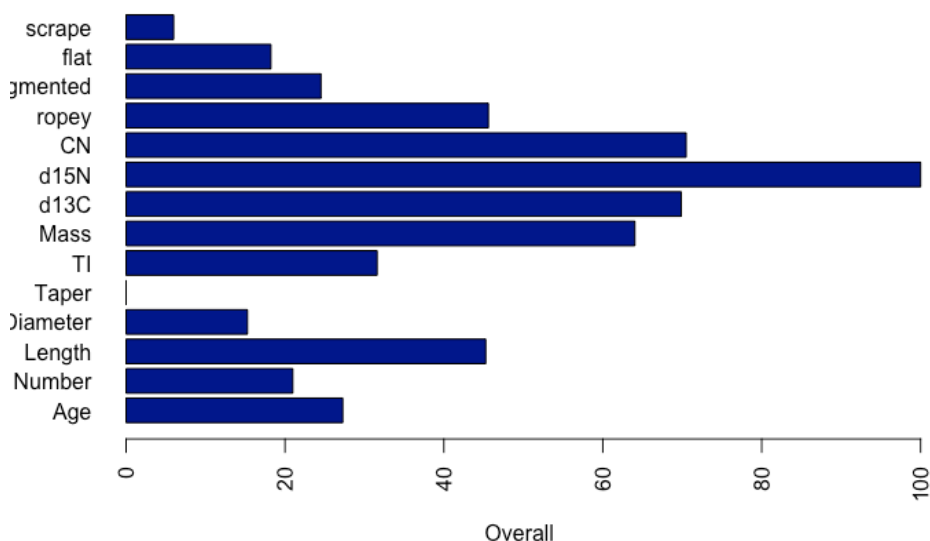
```
var_nnet <- varImp(object=model_nnet)
```

```
barplot(var_nnet$importance[, 'Overall'], main="Neural Networks - Variable Importance",
```

```
       xlab='Overall', horiz=TRUE, las=2,
```

```
       col="darkblue", names.arg=row.names(var_nnet$importance)[])
```

Neural Networks - Variable Importance



```
predictions<-predict.train(object=model_nnet,testSet[,predictors],type="raw")
```

```
confusionMatrix(predictions,testSet[,outcomeName])
```

```
> predictions<-predict.train(object=model_nnet,testSet[,predictors],type="raw")
```

```
> confusionMatrix(predictions,testSet[,outcomeName])
```

Confusion Matrix and Statistics

	Reference		
Prediction	0	1	2
0	14	0	1
1	0	5	1
2	0	2	4

Overall Statistics

Accuracy : 0.8519
95% CI : (0.6627, 0.9581)
No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7551

McNemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.9231	0.9500	0.9048
Pos Pred Value	0.9333	0.8333	0.6667
Neg Pred Value	1.0000	0.9048	0.9048
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.5556	0.2222	0.2222
Balanced Accuracy	0.9615	0.8321	0.7857

```
> |
```

Model Statistics for Naive Bayes

```
print(model_nb)
```

```
> print(model_nb)
```

Naive Bayes

83 samples

14 predictors

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

usekernel	Accuracy	Kappa
FALSE	0.5698618	0.3513484
TRUE	0.6624052	0.4391168

Tuning parameter 'laplace' was held constant at a value of 0

Tuning parameter 'adjust' was

held constant at a value of 1

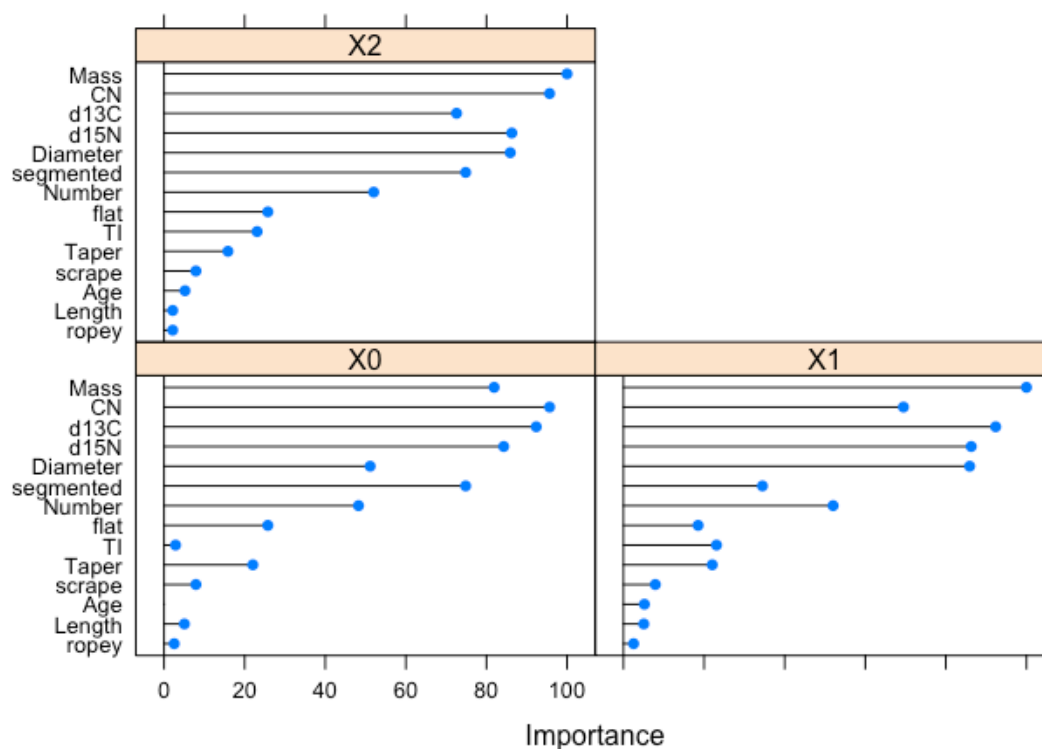
Accuracy was used to select the optimal model using the largest value.

The final values used for the model were laplace = 0, usekernel = TRUE and adjust = 1.

```
> |
```

```
plot(varImp(object=model_nb),main="Naive Bayes - Variable Importance")
```

Naive Bayes - Variable Importance



```
predictions<-predict.train(object=model_nb,testSet[,predictors],type="raw")
```

```
confusionMatrix(predictions,testSet[,outcomeName])
```

```
> predictions<-predict.train(object=model_nb,testSet[,predictors],type="raw")
> confusionMatrix(predictions,testSet[,outcomeName])
```

Confusion Matrix and Statistics

	Reference		
Prediction	0	1	2
0	14	2	2
1	0	5	0
2	0	0	4

Overall Statistics

Accuracy : 0.8519
 95% CI : (0.6627, 0.9581)
 No Information Rate : 0.5185
 P-Value [Acc > NIR] : 0.0003126

Kappa : 0.7416

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.6923	1.0000	1.0000
Pos Pred Value	0.7778	1.0000	1.0000
Neg Pred Value	1.0000	0.9091	0.9130
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.6667	0.1852	0.1481
Balanced Accuracy	0.8462	0.8571	0.8333

```
> |
```


Model Statistics for GBM

```
print(model_gbm)
```

```
> print(model_gbm)
Stochastic Gradient Boosting

83 samples
14 predictors
 3 classes: '0', '1', '2'

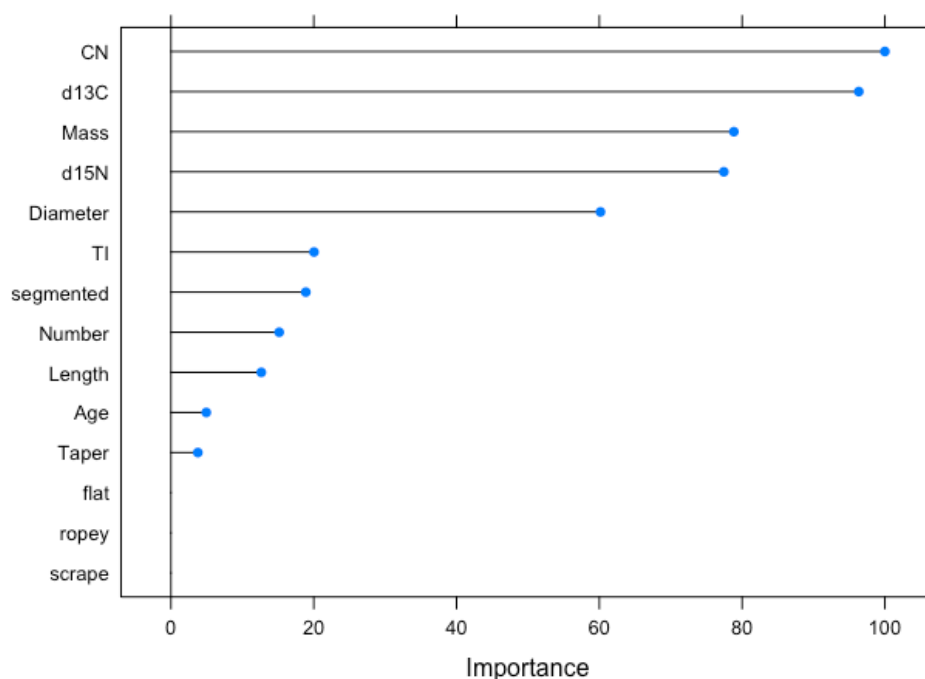
No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
Resampling results across tuning parameters:

 interaction.depth  n.trees  Accuracy  Kappa
1                  50      0.6275335  0.3671868
1                  100      0.6196931  0.3647420
1                  150      0.6142946  0.3550913
2                   50      0.6196146  0.3657719
2                  100      0.6166140  0.3590386
2                  150      0.5936106  0.3231235
3                   50      0.6170225  0.3621061
3                  100      0.6003658  0.3333242
3                  150      0.6027883  0.3373044

Tuning parameter 'shrinkage' was held constant at a value of 0.1
Tuning parameter
'n.minobsinnode' was held constant at a value of 10
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were n.trees = 50, interaction.depth = 1, shrinkage = 0.1
and n.minobsinnode = 10.
> |
```

```
plot(varImp(object=model_gbm),main="GBM - Variable Importance")
```

GBM - Variable Importance



```
predictions<-predict.train(object=model_gbm,testSet[,predictors],type="raw")
confusionMatrix(predictions,testSet[,outcomeName])
```

```
> predictions<-predict.train(object=model_gbm,testSet[,predictors],type="raw")
> confusionMatrix(predictions,testSet[,outcomeName])
Confusion Matrix and Statistics
```

```

      Reference
Prediction 0  1  2
      0 14  1  2
      1  0  5  0
      2  0  1  4

```

Overall Statistics

```

      Accuracy : 0.8519
      95% CI   : (0.6627, 0.9581)
No Information Rate : 0.5185
P-Value [Acc > NIR] : 0.0003126

```

```
      Kappa : 0.7465
```

```
McNemar's Test P-Value : 0.2614641
```

Statistics by Class:

	Class: 0	Class: 1	Class: 2
Sensitivity	1.0000	0.7143	0.6667
Specificity	0.7692	1.0000	0.9524
Pos Pred Value	0.8235	1.0000	0.8000
Neg Pred Value	1.0000	0.9091	0.9091
Prevalence	0.5185	0.2593	0.2222
Detection Rate	0.5185	0.1852	0.1481
Detection Prevalence	0.6296	0.1852	0.1852
Balanced Accuracy	0.8846	0.8571	0.8095

```
> |
```

6

For the BEST performing models of each (randomforest, neural net, naive bayes and gbm)

create and display a data frame that has the following columns:

ExperimentName, accuracy, kappa. Sort the data frame by accuracy. (15 points)

```

model_results <- NULL
model_results <- data.frame("ExperimentName" = "Random
Forrest",model_rf$results[row.names(model_rf$bestTune),c("Accuracy","Kappa")])
newrow <- data.frame("ExperimentName" = "Neural
Networks",model_nnet$results[row.names(model_nnet$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
newrow <- data.frame("ExperimentName" = "Naive
Bayes",model_nb$results[row.names(model_nb$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
newrow <- data.frame("ExperimentName" =
"GBM",model_gbm$results[row.names(model_gbm$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)

```

Printing the results in descending order

```
model_results[order(-model_results$Accuracy),]
```

```
> # Printing the results in descending order
> model_results[order(-model_results$Accuracy),]
  ExperimentName Accuracy      Kappa
6 Neural Networks 0.6908921 0.4788325
2   Naive Bayes 0.6624052 0.4391168
3 Random Forrest 0.6507490 0.4248715
1           GBM 0.6275335 0.3671868
> |
```

7

Tune the GBM model using tune length = 20 and:

a) print the model summary and

b) plot the models. (20 points)

Using tuneLength

```
model_gbm_tuned <- train(trainSet[,predictors],trainSet[,outcomeName],method='gbm',tuneLength=20)
print(model_gbm_tuned)
```

```
> print(model_gbm_tuned)
Stochastic Gradient Boosting

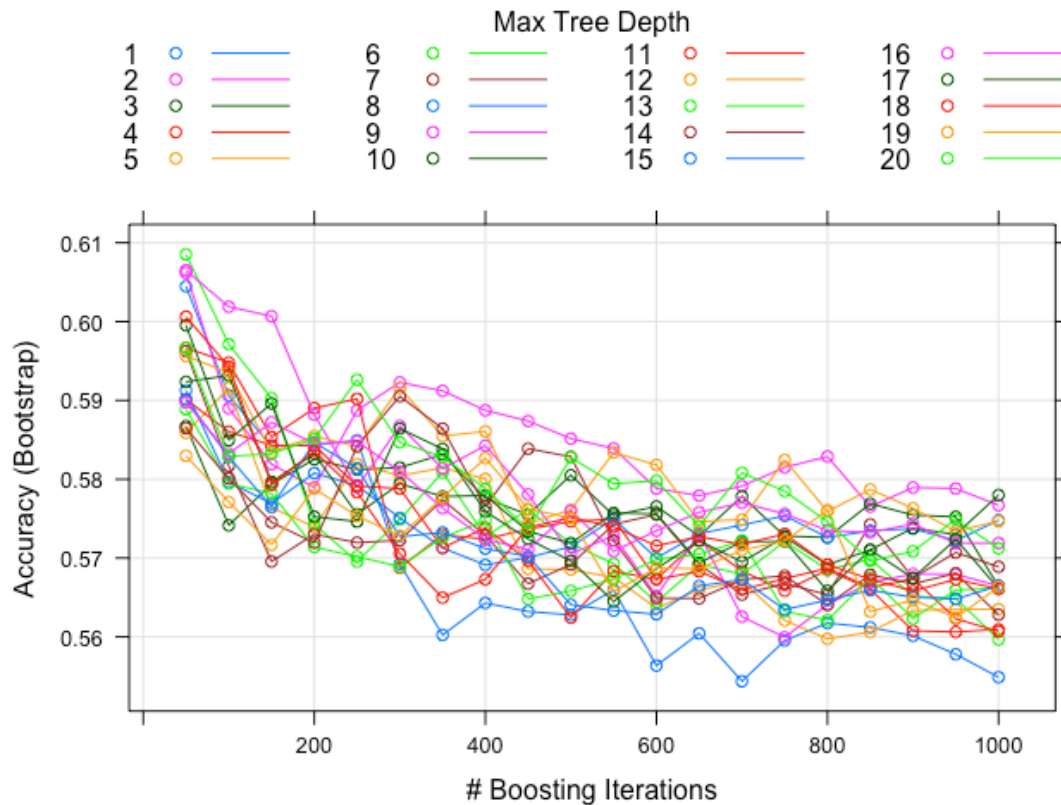
83 samples
14 predictors
 3 classes: '0', '1', '2'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...
Resampling results across tuning parameters:
```

```
Tuning parameter 'shrinkage' was held constant at a value of 0.1
Tuning parameter
  'n.minobsinnode' was held constant at a value of 10
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were n.trees = 50, interaction.depth = 6, shrinkage = 0.1
and n.minobsinnode = 10.
> |
```

visualize the models

```
plot(model_gbm_tuned)
```

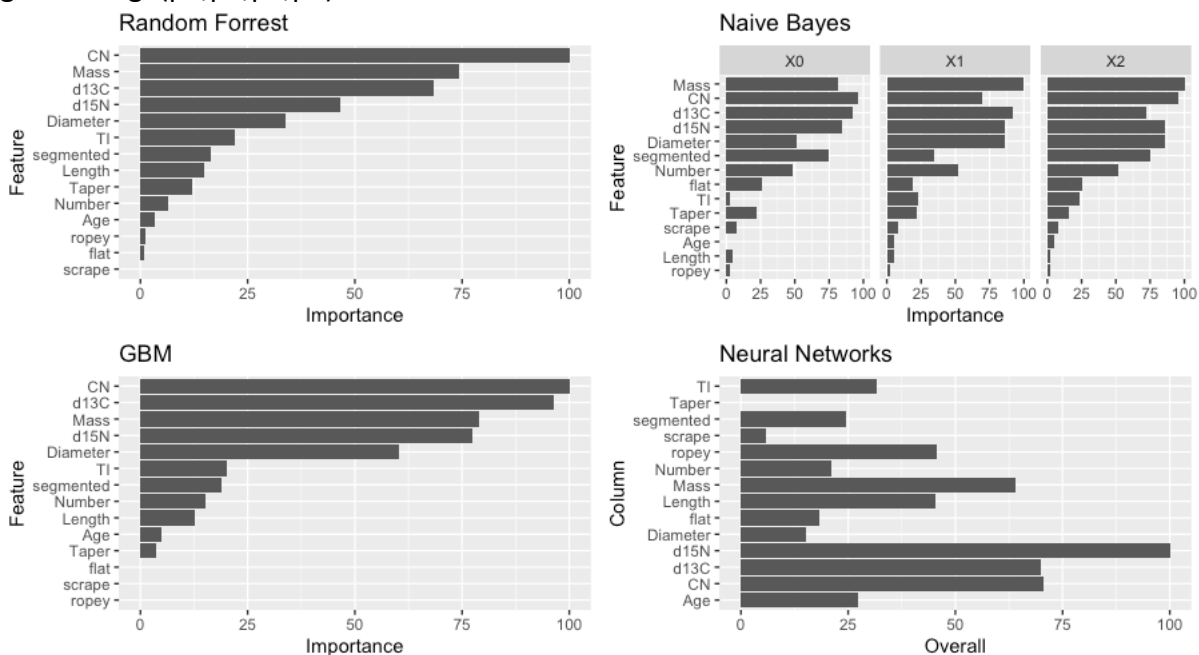


8

Using GGplot and gridExtra to plot all variable of importance plots into one single plot. (10 points)

```
p1 <- ggplot(varImp(object=model_rf)) + ggtitle("Random Forrest")
p2 <- ggplot(varImp(object=model_nb)) + ggtitle("Naive Bayes")
p3 <- ggplot(varImp(object=model_gbm)) + ggtitle("GBM")
var_nnet_df <-
data.frame("Column"=row.names(var_nnet$importance[]),varImp(object=model_nnet)$importance)
p4 <- ggplot(var_nnet_df) + geom_col(aes(x=Column,y=Overall)) + coord_flip() + ggtitle("Neural Networks")
```

grid.arrange(p1,p2,p3,p4)



9

Which model performs the best? and why do you think this is the case? Can we accurately predict species on this dataset? (10 points)

After looking at the accuracies of the models, as of now Neural Network is performing the best.

But since the accuracy is still 69%, I don't think we can predict the species accurately.

It might be because of the low size of data.

10

Graduate Student Questions:

a. Using feature selection with rfe in caret and the repeatedcv method:

Find the top 3 predictors and build the same models as in 6 and 8 with the same parameters. (20 points)

Feature selection using rfe in caret

```
control <- rfeControl(functions = rfFuncs,
  method = "repeatedcv",
  repeats = 3,
  verbose = FALSE)
outcomeName <- 'Species'
predictors <- names(trainSet)[!names(trainSet) %in% outcomeName]
Scat_Pred_Profile <- rfe(trainSet[,predictors], trainSet[,outcomeName], sizes=3, rfeControl = control)
print(Scat_Pred_Profile)
> Scat_Pred_Profile <- rfe(trainSet[,predictors], trainSet[,outcomeName], sizes=3, rfeControl = control)
> print(Scat_Pred_Profile)
```

Recursive feature selection

Outer resampling method: Cross-Validated (10 fold, repeated 3 times)

Resampling performance over subset size:

Variables	Accuracy	Kappa	AccuracySD	KappaSD	Selected
3	0.6988	0.4791	0.1410	0.2530	*
14	0.6888	0.4572	0.1486	0.2655	

The top 3 variables (out of 3):
CN, d13C, d15N

```
> |
```

Taking only the top 3 predictors

```
predictors <- c("CN", "d13C", "d15N")
```

Training Models Using Caret

```
model_rf_fs <- train(trainSet[,predictors], trainSet[,outcomeName], method='rf')
model_nnet_fs <- train(trainSet[,predictors], trainSet[,outcomeName], method='nnet')
model_nb_fs <- train(trainSet[,predictors], trainSet[,outcomeName], method='naive_bayes')
model_gbm_fs <- train(trainSet[,predictors], trainSet[,outcomeName], method='gbm')
```

Tune the GBM model using tune length = 20 and:

a) print the model summary and

b) plot the models.

Using tuneLength

#using tune length

```
model_gbm_tuned_fs <-
```

```
train(trainSet[,predictors],trainSet[,outcomeName],method='gbm',tuneLength=20)
```

```
print(model_gbm_tuned_fs)
```

```
> print(model_gbm_tuned_fs)
```

Stochastic Gradient Boosting

83 samples

3 predictor

3 classes: '0', '1', '2'

No pre-processing

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 83, 83, 83, 83, 83, 83, ...

Resampling results across tuning parameters:

Tuning parameter 'shrinkage' was held constant at a value of 0.1

Tuning parameter

'n.minobsinnode' was held constant at a value of 10

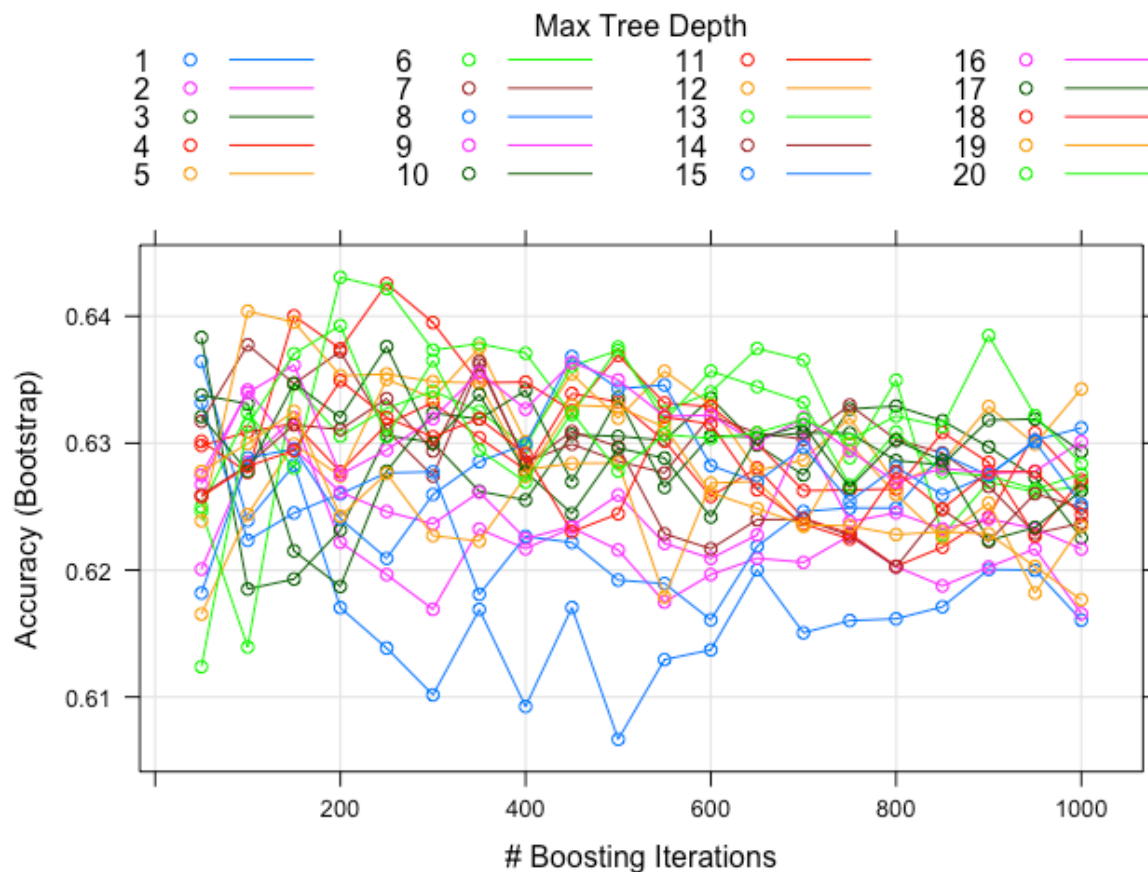
Accuracy was used to select the optimal model using the largest value.

The final values used for the model were n.trees = 200, interaction.depth = 13, shrinkage = 0.1 and n.minobsinnode = 10.

```
> |
```

visualize the models

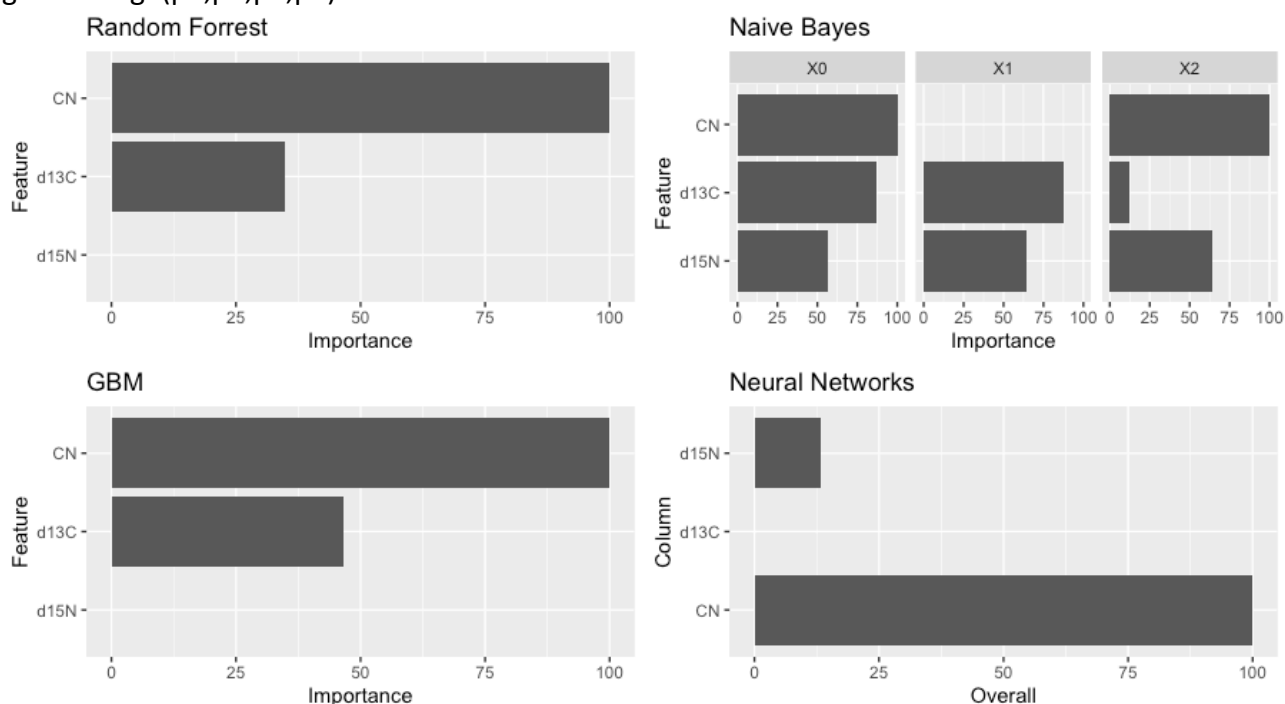
```
plot(model_gbm_tuned_fs)
```



Using GGplot and gridExtra to plot all variable of importance plots into one single plot.

```
p1 <- ggplot(varImp(object=model_rf_fs)) + ggtitle("Random Forrest")
p2 <- ggplot(varImp(object=model_nb_fs)) + ggtitle("Naive Bayes")
p3 <- ggplot(varImp(object=model_gbm_fs)) + ggtitle("GBM")
var_nnet_df_fs <-
data.frame("Column"=row.names(var_nnet_fs$importance)[,varImp(object=model_nnet_fs)$importance)
p4 <- ggplot(var_nnet_df_fs) + geom_col(aes(x=Column,y=Overall)) + coord_flip() + ggtitle("Neural Networks")
```

```
grid.arrange(p1,p2,p3,p4)
```



10

b. Create a dataframe that compares the non-feature selected models (the same as on 7)
 ### and add the best BEST performing models of each (randomforest, neural net, naive bayes and gbm)

and display the data frame that has the following columns: ExperimentName, accuracy, kappa.

Sort the data frame by accuracy.

```
newrow <- data.frame("ExperimentName" = "Random Forrest
FS",model_rf_fs$results[row.names(model_rf_fs$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
newrow = data.frame("ExperimentName" = "Neural Networks
FS",model_nnet_fs$results[row.names(model_nnet_fs$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
newrow = data.frame("ExperimentName" = "Naive Bayes
FS",model_nb_fs$results[row.names(model_nb_fs$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
newrow = data.frame("ExperimentName" = "GBM
FS",model_gbm_fs$results[row.names(model_gbm_fs$bestTune),c("Accuracy","Kappa")])
model_results <- rbind(model_results,newrow)
```

Printing the results in descending order

```

model_results[order(-model_results$Accuracy),]
> # Printing the results in descending order
> model_results[order(-model_results$Accuracy),]
      ExperimentName Accuracy      Kappa
11    Naive Bayes FS 0.7397417 0.5485572
61 Neural Networks FS 0.7204571 0.5313425
6     Neural Networks 0.6908921 0.4788325
21 Random Forrest FS 0.6672625 0.4268536
2      Naive Bayes 0.6624052 0.4391168
3     Random Forrest 0.6507490 0.4248715
7          GBM FS 0.6344563 0.3813186
1          GBM 0.6275335 0.3671868
> |

```

10

c. Which model performs the best?

and why do you think this is the case?

Can we accurately predict species on this dataset? (10 points)

After doing the feature selection, the best model now is Naive Bayes with Feature selection.

Since the number of features has greatly reduced and the target variable has only 3 categories,

It might be the case that Naive Bayes came first.

But even here, the accuracy is only 73%, so we might still not be able to correctly predict species.

With a bigger sample, with the same models we might have better results.